

# Traffic Sign Recognition

Ioannis Ioannidis  
Dimitrios Patiniotis Spyropoulos

University of Piraeus  
NCSR Demokritos  
Athens

27-06-2022

- ▶ Introduction
- ▶ Dataset Description
- ▶ Theoretical Background
- ▶ Experiments
- ▶ Conclusions

## Motivation

- ▶ Road safety is attracting the attention of many researchers around the world.
- ▶ Systems for the detection, classification, and recognition of road signs have become a very important branch on this domain.
- ▶ Traffic sign detectors effectively assist drivers in the process of driving and keep them driving more safely as they inform them about current road situations and potential hazards.
- ▶ Constitute the "eyes" of self-driving cars.



# Dataset Description

## Data Source

- ▶ Dataset for "German Traffic Sign Recognition Benchmark", a multi-category classification competition held at IJCNN 2011.
- ▶ Comprehensive, lifelike dataset of more than 40,000 traffic sign images.
- ▶ Reflects the strong variations in visual appearance of signs due to distance, illumination, weather conditions, partial occlusions, and rotations.






# Theoretical Background

## CNN

- ▶ Convolutional Neural Networks (CNNs) is a state of the art pattern recognition method in computer vision.
- ▶ Unlike traditional neural networks, which work with one-dimensional feature vectors, a CNN takes a two-dimensional image and consequentially processes it with convolutional layers.




A generic CNN Architecture.

# Theoretical Background

## CNN

- ▶ A CNN is composed of input and output layers and multiple hidden layers, which can be divided into a convolution layer, a pooling layer, a rectified linear unit layer, and a fully connected layer.
- ▶ Each convolutional layer consists of a set of trainable filters and computes dot productions between these filters and layer input to obtain an activation map.
- ▶ These filters are also known as kernels and allow detecting the same features in different locations.



## Description

- ▶ Experiment I: Optimizers Comparison
- ▶ Experiment II: Training Size Modification
- ▶ Experiment III: Number of Convolutional Layers





## Optimizers Comparison

### ► Stochastic Gradient Decent

Minimizes an objective function  $J(\theta)$  by updating the parameters in the opposite direction of the gradient of the objective function  $\nabla_{\theta} J(\theta)$  w.r.t. to the parameters.




### ► Adam Optimizer

Accelerates the gradient descent algorithm by taking into consideration the "exponentially weighted average" of the gradients. Using averages makes the algorithm converge towards the minima in a faster pace.

### ► AdaDelta Optimizer


Dynamically adapts over time using only first order information and has minimal computational overhead. It requires no manual tuning of a learning rate and appears robust to noisy gradient information, different model architecture choices, various data modalities and selection of hyperparameters.

## Optimizers Comparison Results




# Experiment II


## Training Size Modification



(a) 5s/epoch



(b) 9s/epoch



(c) 20s/epoch

## Reducing Convolution Layers

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 28, 28, 16)	448
conv2d_1 (Conv2D)	(None, 26, 26, 32)	4640
max_pooling2d (MaxPooling2D)	(None, 13, 13, 32)	0
batch_normalization (Batch Normalization)	(None, 13, 13, 32)	128
conv2d_2 (Conv2D)	(None, 11, 11, 64)	18496
conv2d_3 (Conv2D)	(None, 9, 9, 128)	73856
max_pooling2d_1 (MaxPooling2D)	(None, 4, 4, 128)	0
batch_normalization_1 (Batch Normalization)	(None, 4, 4, 128)	512
flatten (Flatten)	(None, 2048)	0
dense (Dense)	(None, 512)	1049088
batch_normalization_2 (Batch Normalization)	(None, 512)	2048
dropout (Dropout)	(None, 512)	0
dense_1 (Dense)	(None, 43)	22059


=====  
Total params: 1,171,275  
Trainable params: 1,169,931  
Non-trainable params: 1,344




Layer (type)	Output Shape	Param #
conv2d_4 (Conv2D)	(None, 28, 28, 16)	448
max_pooling2d_2 (MaxPooling2D)	(None, 14, 14, 16)	0
batch_normalization_3 (Batch Normalization)	(None, 14, 14, 16)	64
conv2d_5 (Conv2D)	(None, 12, 12, 32)	4640
max_pooling2d_3 (MaxPooling2D)	(None, 6, 6, 32)	0
batch_normalization_4 (Batch Normalization)	(None, 6, 6, 32)	128
flatten_1 (Flatten)	(None, 1152)	0
dense_2 (Dense)	(None, 128)	147584
batch_normalization_5 (Batch Normalization)	(None, 128)	512
dropout_1 (Dropout)	(None, 128)	0
dense_3 (Dense)	(None, 43)	5547

=====  
Total params: 158,923  
Trainable params: 158,571  
Non-trainable params: 352

## Reducing Convolution Layers Results



(a) Baseline CNN



(b) Reduced Architecture

# Conclusion

- ▶ Despite the gradual increase of SGD, Adam seems to perform better and faster, while AdaDelta is way too slow.
- ▶ Concerning training size, we observed that while accuracy seems not to be that much affected, training time is significantly reduced when we decrease the train sample.
- ▶ Finally, reducing the model's convolutions and parameters, we achieve similar performance, with lower computational cost and time.

